

Groupy: a group membership service

Jordi Guitart

Adapted with permission from Johan Montelius (KTH)

February 6, 2025

Introduction

This is an assignment where you will implement a group membership service that provides atomic multicast. The aim is to have several application-layer processes with a **coordinated replicated state**, i.e., they should all perform the same sequence of state changes. A node that wishes to perform a state change must first multicast the change to the group so that all nodes can execute it. Since the multicast layer provides total order, all nodes will be synchronized. The problem is that all nodes need to be synchronized even though nodes may come and go (crash). As you will see it is not as trivial as one might first think.

This assignment includes a number of coding and experimental milestones. On accomplishing each of them, the team must show their outcome to the professor during the practical sessions, so that he can track the progress and evaluate the team accordingly. The team can validate their coding milestones beforehand through the review questionnaires at ATENEA. This assignment also includes a number of open questions. The team must answer each of them through the questionnaire at ATENEA before the deadline (one week after the last session of the seminar).

1 The architecture

We will implement a group membership service that provides atomic multicast in **view synchrony**. The architecture of this service consists of a set of nodes where one is the elected leader. All nodes that wish to multicast a message will send the message to the leader and the leader will do a multicast to all members of the group. If the leader dies a new leader is elected.

A new node that wishes to enter the group will contact any node in the group and request to join. The leader will determine when the node is to be included and will deliver a **new group view** to the group.

The application-layer processes will use this group membership service to synchronize their states. As commented, this service is implemented by means of a set of group processes. Each application-layer process will have its own group process that it communicates with. The application-layer process will send the messages to be multicast to the group process and will receive all multicast messages from it. The group process will tell the application-layer process to deliver a message only when the rest of processes have also delivered it. The application-layer process must also be prepared to decide if a new node should be allowed to join the group and also decide the initial state of this node.

Note that we will not deliver any group views to the application-layer process. We could adapt the system so that it reports any view changes but for the application that we are targeting this is not needed. We will keep it as simple as possible and then discuss extensions and how much they would cost.

1.1 View synchrony

Each node should be able to multicast messages to the members of the group. The communication is divided into views and messages will be said to be delivered in a view. For all messages in a view we will guarantee the following:

- in FIFO order: in the order that they were sent by a node
- in total order: all nodes see the same sequence
- reliably: if a correct node delivers a message, all correct nodes deliver the message

The last statement seems to be a bit weak, what do we mean by a correct node? A node will fail only by crashing and will then never be heard from again. A correct node is a node that does not fail during a view (i.e., it survives to install the next view).

It will not be guaranteed that sent messages are delivered, we will use asynchronous sending without acknowledgment and if we have a failing leader a sent message might disappear.

1.2 The leader

A node will either play the role of a leader (let's hope there is only one) or a slave. The slaves will forward messages to the leader and the leader will tag each message with a sequence number and multicast it to all nodes. The leader can also accept a message directly from its own master (i.e. the application-layer process). The application-layer process is unaware of whether its group process is acting as a leader or a slave.

1.3 A slave

A slave will receive messages from its application-layer process and forward them to the leader. It will also receive messages from the leader and forward them to the application-layer process. If nodes would not fail this would be the easiest job in the world but since we must be able to act if the leader dies we need to do some book-keeping.

In our first version we will not deal with failures but only with adding new nodes to the system. This is complicated enough to start with.

1.4 The election

The election procedure is very simple. All slaves have the same list of peers and they all elect the first node in the list as the leader. A slave that detects that it is the first node will of course adopt the role as leader.

1.5 The application-layer process

An application-layer process will create a group process and contact any other application-layer process it knows of. It will request to join the group providing the process identifier of its group process, which will then wait for a view delivery, containing the peer processes in the group.

There is no guarantee that the request is delivered to the leader, or rather the leader could be dead and we have not detected this yet. The requesting application-layer process is however not told about this so we cannot do anything but wait and hope for the best. We will use a timeout and if we have not been invited in a second we might as well abort the attempt.

Once added to the group, the application-layer process has the problem of obtaining the correct state (the color). It does this by using the atomic multicast system in a clever way. It sends a request to obtain the state and waits for this message to be delivered to itself. It now knows that the other processes have seen this message and will respond by sending the state.

The state message might however not be the first message that is delivered. We might have other state changes in the pipeline. Once the state is received these state changes must of course be applied to the state. The implementation uses the implicit deferral of Erlang and simply lets any state change messages remain in the message queue and chooses to handle the state message first before the state change messages.

2 The first implementation

Our first version, called `gms1`, will only handle starting of a single node and adding more nodes. Failures will not be handled so some of the states that we need to keep track of are not described. We will then extend this implementation to handle failures later on.

A group process will be a slave when started, but might in the future become a leader. The first process that is started will however become a leader directly.

2.1 The leader

The leader keeps the following state:

- **Name:** a unique name of the node, only used for debugging
- **Master:** the process identifier of the application-layer process
- **Slaves:** an ordered list of the process identifiers of all slaves in the group

The list of peers is ordered based on when they were admitted to the group. We will use this order in the election procedure.

The leader should be able to handle the following messages:

- `{mcast, Msg}`: a message either from its own master or from a peer node. A message `{msg, Msg}` is multicasted to all peers (using function `bcast/3` that sends a message to each of the processes in a list) and a message `{deliver, Msg}` is sent to the application-layer process (i.e. **Master**).

- **{join, Peer}**: a message, from a peer or the master, that is a request from a node to join the group. A **view** message containing the new set of slaves is multicasted to all peers.

```

leader(Name, Master, Slaves) ->
  receive
    {mcast, Msg} ->
      bcast(Name, ..., ...), %% TODO: COMPLETE
      %% TODO: ADD SOME CODE
      leader(Name, Master, Slaves);
    {join, Peer} ->
      NewSlaves = lists:append(Slaves, [Peer]),
      bcast(Name, ..., ...), %% TODO: COMPLETE
      leader(Name, Master, ...); %% TODO: COMPLETE
    stop ->
      ok;
    Error ->
      io:format("leader ~s: strange message ~w~n", [Name, Error])
  end.

bcast(_, Msg, Nodes) ->
  lists:foreach(fun(Node) -> Node ! Msg end, Nodes).

```

Notice that we add the new node at the end of the list of slaves. This is important, we want the new node to be the last one to see the view message that we send out. More on this later when we look at failing nodes.

2.2 A slave

A slave has an even simpler job, it will not make any complicated decisions. It is simply forwarding messages from its master to the leader and vice versa. The state of a slave is as follows:

- **Name**: a unique name of the node, only used for debugging
- **Master**: the process identifier of the application-layer process
- **Leader**: the process identifier of the leader
- **Slaves**: an ordered list of the process identifiers of all slaves in the group

The messages from the master and the leader are the following:

- **{mcast, Msg}**: a request from its master to multicast a message, the message is forwarded to the leader.
- **{join, Peer}**: a request from its master to allow a new node to join the group, the message is forwarded to the leader.
- **{msg, Msg}**: a multicasted message from the leader. A message **{deliver, Msg}** is sent to the master.
- **{view, Leader, NewSlaves}**: a multicasted view from the leader. In this version, the slave is only interested in the new set of peers (**NewSlaves**).

```

slave(Name, Master, Leader, Slaves) ->
  receive
    {mcast, Msg} ->
      %% TODO: ADD SOME CODE
      slave(Name, Master, Leader, Slaves);
    {join, Peer} ->
      %% TODO: ADD SOME CODE
      slave(Name, Master, Leader, Slaves);
    {msg, Msg} ->
      %% TODO: ADD SOME CODE
      slave(Name, Master, Leader, Slaves);
    {view, Leader, NewSlaves} ->
      slave(Name, Master, Leader, ...); %% TODO: COMPLETE
  stop ->
  ok;
  Error ->
    io:format("slave ~s: strange message ~w~n", [Name, Error])
end.

```

Since we will not yet deal with failure there is no transition between being a slave and becoming a leader. We will add this later but first let us have this thing up and running.

2.3 Initialization

Initializing a process that is the first node in a group is simple. The only thing we need to do is to give it an empty list of peers. Since it is the only node in the group it will of course be the leader.

```

-module(gms1).
-export([start/1, start/2]).

start(Name) ->
  Self = self(),
  spawn_link(fun()-> init(Name, Self) end).

init(Name, Master) ->
  leader(Name, Master, []).

```

Starting a node that should join an existing group is only slightly more problematic. We need to send a `{join, self()}` message to a node in the group and wait for an invitation. The invitation is delivered as a view message containing everything we need to know. The initial state is of course as a slave.

```

start(Name, Grp) ->
  Self = self(),
  spawn_link(fun()-> init(Name, Grp, Self) end).

init(Name, Grp, Master) ->
  Self = self(),
  Grp ! {join, Self},

```

```

receive
  {view, Leader, Slaves} ->
    Master ! joined,
    slave(Name, Master, Leader, Slaves)
end.

```

2.4 The application-layer process

To do some experiments we create an application-layer process (so-called worker) that uses a gui, which is simply a colored window, to describe its state. The window is initially black or in RGB talk $\{0, 0, 0\}$. This is also the initial state of the worker. Each message that is delivered to the worker is an integer N in some interval, say 1 to 20. A worker will change its state by adding N to the R value and rotate the values. If the state of the worker is $\{R, G, B\}$ and the worker is delivered message N , the new state is $\{G, B, (R+N) \bmod 256\}$.

The color of a worker will thus change over time and the order of messages is of course important. The sequence 5, 12, 2 will of course not create the same color as the sequence 12, 5, 2. If all workers start with a black color and we fail to deliver the messages in the same order the colors of the workers will start to diverge.

A worker and the gui are given at the appendices.

2.5 Experiments

Let's do some testing to see that we can create a group, add some peers, and keep their state coordinated. You can use the following code to start and stop the whole system. Note that we are using the name of the module (i.e. `gms1`) as the parameter `Module` to the start procedure. All the workers but the first one need to know a member of the group in order to join. `Sleep` stands for up to how many milliseconds the workers should wait until the next message is sent.

```

-module(groupy).
-export([start/2, stop/0]).

start(Module, Sleep) ->
  register(a, worker:start("P1", Module, Sleep)),
  register(b, worker:start("P2", Module, a, Sleep)),
  register(c, worker:start("P3", Module, b, Sleep)),
  register(d, worker:start("P4", Module, c, Sleep)),
  register(e, worker:start("P5", Module, d, Sleep)).

stop() ->
  stop(a),
  stop(b),
  stop(c),
  stop(d),
  stop(e).

stop(Name) ->
  case whereis(Name) of

```

```

undefined ->
    ok;
Pid ->
    Pid ! stop
end.

```

Note that the test procedure runs all the processes in the same Erlang instance. To run the tests in a distributed fashion, you must also adapt the `groupy` module to start (and stop) each worker in a different Erlang instance. Check the slides about Erlang to refresh how processes are created remotely, how names registered in remote nodes are referred, and how Erlang runtime should be started to run distributed programs.

CODE REVIEW (2 points)	
Leader-slave group membership service: Erlang file <code>gms1.erl</code> .	
Distributed execution: Erlang file <code>groupy.erl</code> .	

EXPERIMENTAL MILESTONES (1.25 points)	
MS1 (0.5 points)	Make tests with different <code>Sleep</code> values and check if the peers can keep their state coordinated.
MS2 (0.25 points)	Crash the leader worker (by sending a <code>stop</code> message from the Erlang prompt) and see what happens.
MS3 (0.5 points)	Try the group membership service while running each worker in a different Erlang instance.

OPEN QUESTIONS (0.25 points)	
a)	What happens if the leader worker crashes?

3 Handling failures

We will build up our fault tolerance gradually. First we will make sure that we detect crashes, then ensure that a new leader is elected and finally make sure that the multicast service preserves the properties of the atomic multicast. Keep `gms1` as a reference and call the adapted module `gms2`.

3.1 Failure detectors

We will use the Erlang built-in support to detect and report that processes have crashed. A process can monitor another process and if the latter dies a message will be received. For now, we will assume that the monitors are perfect, i.e., they will eventually report the crash of a process and they will never report the death of a process that has not died.

We will also assume that the message that informs a process about a death of another process is the last message that it will see from it. The message will thus be received in FIFO order as any regular message.

The question we first need to answer is, who should monitor who? We will keep things simple: the only node that will be monitored is the leader. We will not report new views when a slave dies.

A slave that detects that the leader has died has to move to an election state. This is implemented by adding a call to `erlang:monitor/2` in the initialization of the slave to enable the monitoring of the leader and a new clause in the state of the slave to react to the crash of the leader:

```
{'DOWN', _Ref, process, Leader, _Reason} ->
    election(..., ..., ...); %% TODO: COMPLETE
```

Additionally, since now the leader can crash it could be that a node that wants to join the group will never receive a reply. Its message could be forwarded to a dead leader and the joining node is never informed of the fact that its request was lost. We address this by adding a timeout (for instance, 1000 ms) when waiting for an invitation to join the group in the initialization of the slave.

```
after ?timeout ->
    Master ! {error, "no reply from leader"}
```

The election procedure will select the first process in the list of peers as the leader. If a process finds itself being the first node in the list, it will thus become the leader of the group and **multicast a new view to the rest**. If a process finds that it must remain as a slave, it will start monitoring the new elected leader.

```
election(Name, Master, Slaves) ->
    Self = self(),
    case Slaves of
        [Self|Rest] ->
            %% TODO: ADD SOME CODE
            leader(..., ..., ...); %% TODO: COMPLETE
        [NewLeader|Rest] ->
            NewRef = ... %% TODO: COMPLETE
            slave(..., ..., ..., ..., ...) %% TODO: COMPLETE
    end.
```

One thing that we have to pay attention to is what we should do if, as a slave, we receive the **view** message from the new leader before we have noticed that the old leader is dead. In that case, we should accept the new view, enable the monitoring of the new leader and ignore trailing DOWN messages with the reference of the old leader. Note that **you need to extend the slave method with the reference of the process that we are monitoring**, which must be updated accordingly when we change the monitored process.

```
{view, NewLeader, NewSlaves} ->
    erlang:demonitor(Ref, [flush]),
    NewRef = ... %% TODO: COMPLETE
    slave(Name, Master, ..., ..., ...); %% TODO: COMPLETE
```


Of course, if we receive a view message and the leader has not changed, we should just accept the new view, without changing the monitored process, as we were doing before.

CODE REVIEW (1 points)	
Failure-aware group membership service: Erlang file <code>gms2.erl</code> .	

EXPERIMENTAL MILESTONES (0.5 points)	
MS4 (0.5 points)	Repeat previous tests and check if the peers can keep their state coordinated even if leader and/or slave workers crash.

OPEN QUESTIONS (0.5 points)	
b) What happens now if the leader worker crashes? c) How can the membership service tolerate the crash of a slave worker if we are not monitoring slaves?	

3.2 Missing messages

It seems to be too easy and unfortunately it is not. To show that current implementation is not working we can change the `bcast/3` procedure and introduce a random crash. We define a constant `arghh` that defines the risk of crashing. A value of 100 means that a process will crash in average once in a hundred attempts. The definition of `bcast/3` now looks like this:

```
bcast(Name, Msg, Nodes) ->
    lists:foreach(fun(Node) ->
                    Node ! Msg,
                    crash(Name, Msg)
                  end,
                  Nodes).

crash(Name, Msg) ->
    case rand:uniform(?arghh) of
        ?arghh ->
            io:format("leader ~s CRASHED: msg ~w~n", [Name, Msg]),
            exit(no_luck);
        _ ->
            ok
    end.
```

EXPERIMENTAL MILESTONES (0.5 points)	
MS5 (0.5 points)	Repeat again the tests and see if you can make the state of the workers become out of sync.

OPEN QUESTIONS (0.25 points)
d) Why do the workers desynchronize?

3.3 Reliable multicast

To remedy the problem, we will keep a copy of the last message that we have seen from the leader. If we detect the death of the leader it could be that it died during the multicast procedure and that some nodes have not seen the message. Let's make an assumption that we will discuss later:

- Messages are reliably delivered and thus, if the leader sends a message to slave A and then to slave B, and B receives the message, then also A will receive the message.

The leader is sending messages to the peers in the order that they occur in the list of peers. If anyone receives a message then the first peer in the list has also received the message. This means that the next leader can resend that message to the rest.

This will of course introduce the possibility of doublets of messages being received. In order to detect this we will number all messages and only deliver new messages to the application-layer process.

Let's go through the changes that we need to make. Create a new module `gms3` that implements them.

- The `slave` procedure is extended with two arguments: `N` and `Last`. `N` is the expected sequence number of the next message and `Last` is a copy of the last message (**a regular multicast message or a view**) received from the leader.
- The `election` procedure is extended with the same two arguments.
- The `leader` procedure is extended with the argument `N`, the sequence number of the next message (regular message or view) to be sent.
- The multicast messages (`msg` and `view`) also change and will now contain the sequence number.
- We must also add clauses to the slave to accept and **ignore duplicate msg and view messages** from the leader. If we do not remove these from the message queue they will add up and after a year generate a very hard to handle trouble report. When discarding messages we only want to discard messages that we have seen, i.e., messages with a sequence number less than `N`. We can do it by using the `when` construction. For example:

`{msg, _, I} when I < N -> ...`

- The crucial part is then in the election procedure where **the elected leader will multicast the last received message to all peers** in the group before changing the view.

CODE REVIEW (2.5 points)	
Reliable group membership service: Erlang file <code>gms3.erl</code> .	

EXPERIMENTAL MILESTONES (0.5 points)	
MS6 (0.5 points)	Repeat the experiments to see if now the peers can keep their state coordinated even if nodes crash.

Assuming all tests went well we're ready to ship the product. There is however one thing we need to mention and that is that our implementation does not work!!! Well, it sort of works depending on what the Erlang environment guarantees and how strong our requirements are.

3.4 What could possibly go wrong

The first thing we have to realize is what guarantees the Erlang system actually gives on message sending: it only guarantees that messages are delivered in FIFO order, not that they actually do arrive. We have built our system relying on reliable delivery of messages, something that is not guaranteed.

The second reason why things will not work is that we rely on that the Erlang failure detector is perfect, i.e., that it will never suspect any correct node for having crashed, and implementing a perfect failure detector in purely asynchronous systems is impossible.

The third reason why things do not work is that we could have a situation where one node that will eventually crash delivers a message that will not be delivered by any correct node. This could happen even if we had reliable send operations and perfect failure detectors, and to be avoided requires a multicast implementation providing uniform agreement (which allows a process to crash after it has delivered a message, while still ensuring that all correct processes will deliver the message).

OPEN QUESTIONS (0.75 points)	
e)	How should we change the implementation to handle the potential loss of messages in the network? How would this change impact performance?
f)	What would happen if we wrongly suspect the leader to have crashed?
g)	When is it possible that a process crashes after it has delivered a message that will not be delivered by any correct node?

Appendix A: *worker.erl*

```
-module(worker).
-export([start/3, start/4]).

-define(change, 20).
-define(color, {0,0,0}).

start(Name, Module, Sleep) ->
    spawn(fun() -> init(Name, Module, Sleep) end).

init(Name, Module, Sleep) ->
    Cast = apply(Module, start, [Name]),
    Color = ?color,
    init_cont(Name, Cast, Color, Sleep).

start(Name, Module, Peer, Sleep) ->
    spawn(fun() -> init(Name, Module, Peer, Sleep) end).

init(Name, Module, Peer, Sleep) ->
    Cast = apply(Module, start, [Name, Peer]),
    receive
        joined ->
            Ref = make_ref(),
            Cast ! {mcast, {state_req, Ref}},
            Color = state_transfer(Cast, Ref),
            if Color /= stop ->
                init_cont(Name, Cast, Color, Sleep),
                Cast ! stop;
            true ->
                Cast ! stop
            end;
        {error, Error} ->
            io:format("worker ~s: error: ~s~n", [Name, Error]);
        stop ->
            ok
    end.

state_transfer(Cast, Ref) ->
    receive
        {deliver, {state_req, Ref}} ->
            receive
                {deliver, {set_state, Ref, Color}} ->
                    Color
            end;
        {join, Peer} ->
            Cast ! {join, Peer},
            state_transfer(Cast, Ref);
        stop ->
            stop;
    end.
```

```

        _Ignore ->
            state_transfer(Cast, Ref)
    end.

init_cont(Name, Cast, Color, Sleep) ->
    Gui = gui:start(Name, self()),
    Gui ! {color, Color},
    Wait = if Sleep == 0 -> 0; true -> rand:uniform(Sleep) end,
    timer:send_after(Wait, cast_change),
    worker(Name, Cast, Color, Gui, Sleep),
    Gui ! stop.

worker(Name, Cast, Color, Gui, Sleep) ->
    receive
        {deliver, {change_state, N}} ->
            NewColor = change_color(N, Color),
            Gui ! {color, NewColor},
            worker(Name, Cast, NewColor, Gui, Sleep);
        {deliver, {state_req, Ref}} ->
            Cast ! {mcast, {set_state, Ref, Color}},
            worker(Name, Cast, Color, Gui, Sleep);
        {deliver, {set_state, _, _}} ->
            worker(Name, Cast, Color, Gui, Sleep);
        {join, Peer} ->
            Cast ! {join, Peer},
            worker(Name, Cast, Color, Gui, Sleep);
        cast_change ->
            Cast ! {mcast, {change_state, rand:uniform(?change)}},
            Wait = if Sleep == 0 -> 0; true -> rand:uniform(Sleep) end,
            timer:send_after(Wait, cast_change),
            worker(Name, Cast, Color, Gui, Sleep);
        stop ->
            Cast ! stop,
            ok;
        Error ->
            io:format("worker ~s: strange message: ~w~n", [Name, Error]),
            worker(Name, Cast, Color, Gui, Sleep)
    end.

change_color(N, {R,G,B}) ->
    {G, B, ((R+N) rem 256)}.

```

Appendix B: *gui.erl*

```
-module(gui).
-export([start/2]).
-define(width, 200).
-define(height, 200).
-include_lib("wx/include/wx.hrl").

start(Name, Master) ->
    spawn_link(fun() -> init(Name, Master) end).

init(Name, Master) ->
    Frame = make_frame(Name),
    loop(Frame, Master).

make_frame(Name) ->          %Name is the window title
    Server = wx:new(), %Server will be the parent for the Frame
    Frame = wxFrame:new(Server, -1, Name, [{size,{?width, ?height}}]),
    wxFrame:setBackgroundColour(Frame, ?wxBLACK),
    wxFrame:show(Frame),
    %monitor closing window event
    wxFrame:connect(Frame, close_window),
    Frame.

loop(Frame, Master)->
    receive
        %check if the window was closed by the user
        #wx{event=#wxClose{}} ->
            wxWindow:destroy(Frame),
            Master ! stop,
            ok;
        {color, Color} ->
            color(Frame, Color),
            loop(Frame, Master);
        stop ->
            ok;
        Error ->
            io:format("gui: strange message ~w ~n", [Error]),
            loop(Frame, Master)
    end.

color(Frame, Color) ->
    wxFrame:setBackgroundColour(Frame, Color),
    wxFrame:refresh(Frame).
```